

# Parallel Algorithms

## Single-Source Shortest Path

Computação de Alto Desempenho 2024/25

Hervé Paulino

Slides adapted from Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar,  
“Introduction to Parallel Computing”, Addison Wesley, 2003

# Bibliography



---

- Chapter 10 (10.1 to 10.3 and 10.7.2) of Introduction to Parallel Computing (2nd Edition) 2nd Edition, Ananth Grama, George Karypis, Vipin Kumar and Anshul Gupta. Pearson, 2003

# Graph Algorithms

- Graph theory important in computer science
- Many complex problems are graph problems

- $G = (V, E)$ 
  - $V$  finite set of points called vertices
  - $E$  finite set of edges
  - $e \in E$  is an pair  $(u,v)$ , where  $u,v \in V$
  - Unordered and ordered graphs




# Graph Terminology

- Vertex adjacency if  $(u,v)$  is an edge
- Path from  $u$  to  $v$  if there is an edge sequence starting at  $u$  and ending at  $v$
- If there exists a path,  $v$  is reachable from  $u$
- A graph is connected if all pairs of vertices are connected by a path
- A weighted graph associates weights with each edge
- Adjacency matrix is an  $n \times n$  array  $A$  such that
  - $A_{i,j} = 1$  if  $(v_i, v_j) \in E$ ; 0 otherwise
  - Can be modified for weighted graphs ( $\infty$  is no edge)
  - Can be represented as adjacency lists



# Graph Representations

- Adjacency matrix



$$A = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$


- Adjacency list



# Graph Representations

## Compressed Sparse Row (CSR) format

- Represents a matrix with 3 one-dimensional arrays
  - Data: non zero entries of the matrix
  - Columns: the column of each value
  - Rows: the index where each row starts (may contain final element with the number of non-zero elements)
- Used in many GPU Graph processing frameworks
  - Limitations:
    - expensive to access the neighbors of a node
    - Hard to add and delete nodes
  - Some frameworks use more complex methods



$$A = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$

Data	1	1	1	1	1	1	1	1	1	1
columns	1	0	2	4	1	4	4	1	2	3
rows	0	1	4	6	7	10				

# Graph Partitioning: Challenges

- Localizing portions of the computation
  - How to partition the workload so that nearby nodes in the graph are mapped to the same processor?
  - How to partition the workload so that edges that represent significant communication are co-located on the same processor?
- Balancing the load
  - How to give each processor a comparable amount of work?
  - How much knowledge of the graph do we need to do this since complete traversal may not be realistic?
- All of these issues must be addressed by a graph partitioning algorithm that maps individual subgraphs to individual processors.




Image taken from <https://spark.apache.org/docs/latest/graphx-programming-guide.html#pregel-api>

# Graph Processing Models

---

- Vertex centric
  - The kernel is applied over a set of vertexes
- Edge centric
  - The kernel is applied over a set of edges
- There is no clean better solution, it depends on the kernel to apply
  - Frameworks tend to favor vertex-centric models



# Big Data Graph Processing Models

- Pregel
  - Developed by Google
  - Based on the BSP model
  - Vertex centric
  - A graph is divided into partitions: set of vertices plus outgoing edges
  - An hash function on the vertexes id defines to which partition is it assigned
- GraphX
  - Uses the Pregel API
  - Vertex and Edge RDDs
  - Optimized representation
    - Split along vertexes

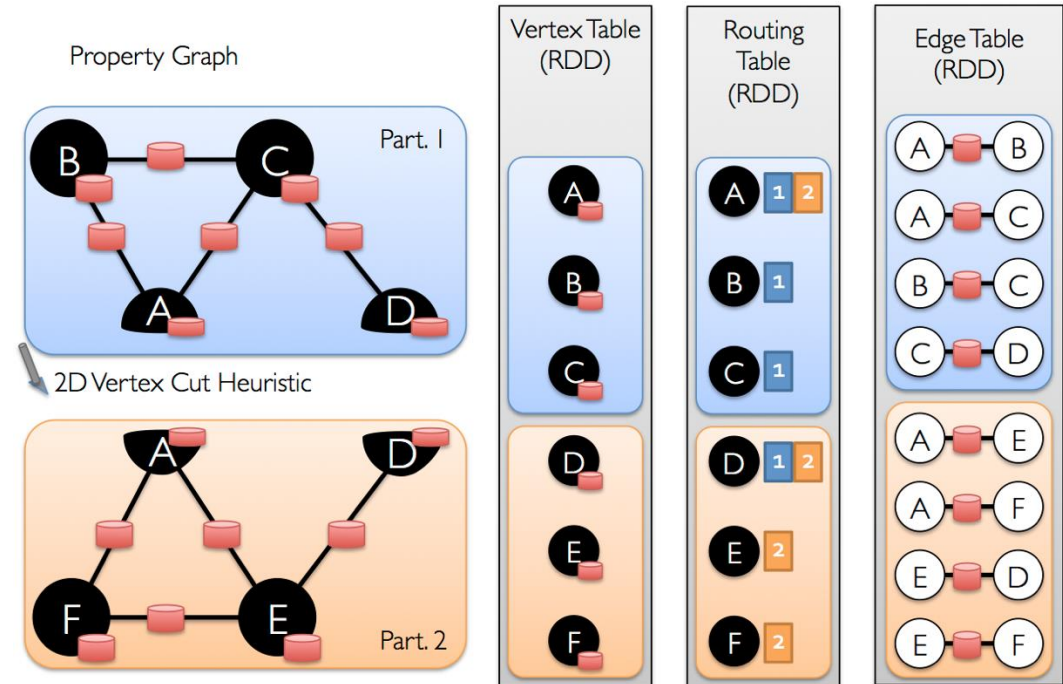


Image taken from <https://spark.apache.org/docs/latest/graphx-programming-guide.html#pregel-api>

# Widely Used Graph Algorithms

- Minimum Spanning Tree
  - A spanning tree of an undirected graph  $G$  is a subgraph of  $G$  that is a tree containing all the vertices of  $G$
  - The minimum spanning tree (MST) for a weighted undirected graph is a spanning tree with minimum weight
- Single Source Shortest Path (SSSP)
  - Shortest paths from a node to all the others
- All Sources Shortest Path
  - Shortest paths between all nodes
- Page Rank
  - Node ranking base on their in-degree and the rank of the nodes that point to it
- Maximal Independent Set
  - A set of vertices  $I \subset V$  is called independent if no pair of vertices in  $I$  is connected via an edge in  $G$ .

# Single-Source Shortest Path

---

- Find shortest path from a vertex  $v$  to all other vertices
- The shortest path in a weighted graph is the edge with the minimum weight
- Weights may represent time, cost, loss, or any other quantity that accumulates additively along a path
- Dijkstra's algorithm finds shortest paths from vertex  $s$ 
  - Incrementally finds shortest paths in greedy manner
  - Keep track of minimum cost to reach a vertex from  $s$
  - $O(n^2)$

# Dijkstra's Single-Source Shortest Path

DIJKSTRA SINGLE SOURCE SP( $V, E, w, s$ ) {

$VT := \{s\};$

**for all**  $v \in (V - VT)$

**if**  $(s, v)$  exists set  $l[v] := w(s, v);$

**else** set  $l[v] := \infty;$

**while**  $VT \neq V$

    find a vertex  $u$  such that  $l[u] := \min\{l[v] \mid v \in (V - VT)\};$

$VT := VT \cup \{u\};$


**for all**  $v \in (V - VT)$

$l[v] := \min\{l[v], l[u] + w(u, v)\};$

}

# SSSP for Vertex A

## Graph




	A 1	B 2	C 3	D 4	E 5
L					

	A 1	B 2	C 3	D 4	E 5
A 1	0	6	∞	6	7
B 2	∞	0	5	∞	∞
C 3	∞	∞	0	9	3
D 4	∞	∞	9	0	7
E 5	∞	4	∞	∞	0

# SSSP for Vertex A

## Graph




	A 1	B 2	C 3	D 4	E 5
L	0	6	$\infty$	6	7

	A 1	B 2	C 3	D 4	E 5
A 1	0	6	$\infty$	6	7
B 2	$\infty$	0	5	$\infty$	$\infty$
C 3	$\infty$	$\infty$	0	9	3
D 4	$\infty$	$\infty$	9	0	7
E 5	$\infty$	4	$\infty$	$\infty$	0

# SSSP for Vertex A

## Graph




	A 1	B 2	C 3	D 4	E 5
L	0	6	11	6	7

	A 1	B 2	C 3	D 4	E 5
A 1	0	6	$\infty$	6	7
B 2	$\infty$	0	5	$\infty$	$\infty$
C 3	$\infty$	$\infty$	0	9	3
D 4	$\infty$	$\infty$	9	0	7
E 5	$\infty$	4	$\infty$	$\infty$	0

# SSSP for Vertex A

## Graph




	A 1	B 2	C 3	D 4	E 5
L	0	6	11	6	7

	A 1	B 2	C 3	D 4	E 5
A 1	0	6	$\infty$	6	7
B 2	$\infty$	0	5	$\infty$	$\infty$
C 3	$\infty$	$\infty$	0	9	3
D 4	$\infty$	$\infty$	9	0	7
E 5	$\infty$	4	$\infty$	$\infty$	0



# SSSP for Vertex A

## Graph




	A 1	B 2	C 3	D 4	E 5
L	0	6	11	6	7

	A 1	B 2	C 3	D 4	E 5
A 1	0	6	$\infty$	6	7
B 2	$\infty$	0	5	$\infty$	$\infty$
C 3	$\infty$	$\infty$	0	9	3
D 4	$\infty$	$\infty$	9	0	7
E 5	$\infty$	4	$\infty$	$\infty$	0

# SSSP for Vertex A

## Graph



	A 1	B 2	C 3	D 4	E 5
L	0	6	11	6	7

	A 1	B 2	C 3	D 4	E 5
A 1	0	6	$\infty$	6	7
B 2	$\infty$	0	5	$\infty$	$\infty$
C 3	$\infty$	$\infty$	0	9	3
D 4	$\infty$	$\infty$	9	0	7
E 5	$\infty$	4	$\infty$	$\infty$	0

# Parallel Formulation of Dijkstra's Algorithm

- Difficult to perform different iterations of the while loop in parallel because  $l[v]$  may change each time
- Classic parallel algorithm paralyzes each iteration:
  - Partition vertices into  $p$  subsets  $V_i$ ,  $i=0, \dots, p-1$
  - Each process  $P_i$  computes  $l_i[u] = \min\{l_i[v] \mid v \in (V - V_T) \cap V_i\}$
  - Global minimum is obtained using all-to-one reduction
  - New vertex is added to  $V_T$  and broadcast to all processes
  - New values of  $l[v]$  are computed for local vertex
- Is this algorithm implementable in GPUs or Spark?

No. Because it is not possible to implement grid-wide **all-one-reduction** reduction, nor **broadcasts**.  
**Adaptations must be made**

# Single-Source Shortest Path

- Dijkstra's algorithm, modified to handle sparse graphs is called Johnson's algorithm.
- The modification accounts for the fact that the minimization step in Dijkstra's algorithm needs to be performed only for those nodes adjacent to the previously selected nodes (the frontier).
- Johnson's algorithm uses a priority queue  $Q$  to store the value  $l[v]$  for each vertex  $v \in (V - V_T)$ .
  - Smaller values (shortest paths) have higher priority


# Single-Source Shortest Path: Parallel Johnson's Algorithm

---

- Maintaining strict order of Johnson's algorithm generally leads to a very restrictive class of parallel algorithms.
- We need to allow exploration of multiple nodes concurrently. This is done by simultaneously extracting  $p$  nodes from the priority queue, updating the neighbors' cost, and augmenting the shortest path.
- If an error is made, it can be discovered (as a shorter path) and the node can be reinserted with this shorter path.

# Single-Source Shortest Paths: Parallel Johnson's Algorithm

- An example of the modified Johnson's algorithm for processing unsafe vertices concurrently.



Priority Queue

(1)  $b:1, d:7, c:\text{inf}, e:\text{inf}, f:\text{inf}, g:\text{inf}, h:\text{inf}, i:\text{inf}$


Array  $l[]$

$a$	$b$	$c$	$d$	$e$	$f$	$g$	$h$	$i$
0	1	$\infty$	7	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$

Shortest path from vertex  $a$

# Single-Source Shortest Paths: Parallel Johnson's Algorithm

- An example of the modified Johnson's algorithm for processing unsafe vertices concurrently.



Shortest path from vertex a

Priority Queue

(1)  $b:1, d:7$ ,  $c:\text{inf}, e:\text{inf}, f:\text{inf}, g:\text{inf}, h:\text{inf}, i:\text{inf}$

(2)  $e:3, c:4, g:10, f:\text{inf}, h:\text{inf}, i:\text{inf}$


Array l[]

a	b	c	d	e	f	g	h	i
0	1	$\infty$	7	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$

a	b	c	d	e	f	g	h	i
0	1	4	7	3	$\infty$	10	$\infty$	$\infty$

# Single-Source Shortest Paths: Parallel Johnson's Algorithm

- An example of the modified Johnson's algorithm for processing unsafe vertices concurrently.



Priority Queue

(1)  $b:1, d:7, c:\text{inf}, e:\text{inf}, f:\text{inf}, g:\text{inf}, h:\text{inf}, i:\text{inf}$

(2)  $e:3, c:4, g:10, f:\text{inf}, h:\text{inf}, i:\text{inf}$

(3)  $h:4, f:6, i:\text{inf}$

Array  $l[]$

$a$	$b$	$c$	$d$	$e$	$f$	$g$	$h$	$i$
0	1	$\infty$	7	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$

0	1	4	7	3	$\infty$	10	$\infty$	$\infty$
---	---	---	---	---	----------	----	----------	----------


0	1	4	7	3	6	10	4	$\infty$
---	---	---	---	---	---	----	---	----------

Shortest path from vertex a



# Single-Source Shortest Paths: Parallel Johnson's Algorithm

- An example of the modified Johnson's algorithm for processing unsafe vertices concurrently.



Priority Queue

- (1)  $b:1, d:7, c:\text{inf}, e:\text{inf}, f:\text{inf}, g:\text{inf}, h:\text{inf}, i:\text{inf}$
- (2)  $e:3, c:4, g:10, f:\text{inf}, h:\text{inf}, i:\text{inf}$
- (3)  $h:4, f:6, i:\text{inf}$
- (4)  $g:5, i:6$

Array l[]

<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>	<i>g</i>	<i>h</i>	<i>i</i>
0	1	$\infty$	7	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
0	1	4	7	3	$\infty$	10	$\infty$	$\infty$
0	1	4	7	3	6	10	4	$\infty$
0	1	4	7	3	6	5	4	6

Shortest path from vertex a

# Single-Source Shortest Paths: Parallel Johnson's Algorithm

- Even if we can extract and process multiple nodes from the queue, the queue itself is a major bottleneck.
- For this reason, we use multiple queues, one for each processor. Each processor builds its priority queue only using its own vertices.
- When process  $P_i$  extracts the vertex  $u \in V_i$ , it sends a message to processes that store vertices adjacent to  $u$ .
- Process  $P_j$ , upon receiving this message, sets the value of  $l[v]$  stored in its priority queue to  $\min\{l[v], l[u] + w(u, v)\}$ .

Once again, in GPUs and Spark this communication requires communication between workers

# Single-Source Shortest Paths: Parallel Johnson's Algorithm

---

- If a shorter path has been discovered to node  $v$ , it is reinserted back into the local priority queue.
- The algorithm terminates only when all the queues become empty.
- A number of node partitioning schemes can be used to exploit graph structure for performance.